

## Project -1

```
from sklearn.linear_model import LogisticRegression
from sklearn.ensemble import RandomForestClassifier
from xgboost import XGBClassifier
from sklearn.metrics import (
    accuracy_score, precision_score, recall_score, f1_score,
    roc_auc_score,
    roc_curve, auc, confusion_matrix, classification_report
)
from sklearn.pipeline import Pipeline
import matplotlib.pyplot as plt
import seaborn as sns
import pandas as pd
from sklearn.pipeline import Pipeline
from sklearn.compose import ColumnTransformer
from sklearn.preprocessing import StandardScaler, OneHotEncoder
from sklearn.impute import SimpleImputer
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
from sklearn.ensemble import RandomForestClassifier
from xgboost import XGBClassifier
from sklearn.metrics import accuracy_score, precision_score,
recall_score, f1_score, roc_auc_score
from sklearn.pipeline import Pipeline

path = "/content/Loan_default.csv"
df = pd.read_csv(path)
df.head()

df.columns
target = "Default"
X = df.drop(columns=[target])
y = df[target]
numeric_cols = X.select_dtypes(include=['int64',
'float64']).columns.tolist()
categorical_cols = X.select_dtypes(include=['object']).columns.tolist()

print("Numeric:", numeric_cols)
print("Categorical:", categorical_cols)
X["Income_to_Loan"] = X["Income"] / (X["LoanAmount"] + 1)
```

```

X["CreditLines_per_Year"] = X["NumCreditLines"] / ((X["MonthsEmployed"] / 12) + 1)

Numeric_cols
=X.select_dtypes(include=['int64','float64']).columns.tolist()
numeric_transformer = Pipeline(steps=[
    ("imputer", SimpleImputer(strategy="median")),
    ("scaler", StandardScaler())
])

categorical_transformer = Pipeline(steps=[
    ("imputer", SimpleImputer(strategy="most_frequent")),
    ("encoder", OneHotEncoder(handle_unknown="ignore"))
])

preprocessor = ColumnTransformer(
    transformers=[
        ("num", numeric_transformer, numeric_cols),
        ("cat", categorical_transformer, categorical_cols)
    ]
)

x_train, x_test, y_train, y_test = train_test_split(
    X, y,
    test_size=0.2,
    random_state=42,
    stratify=y
)

models = {
    "Logistic Regression": LogisticRegression(max_iter=200),

    "Random Forest": RandomForestClassifier(
        n_estimators=100,
        max_depth=10,
        n_jobs=-1,
        random_state=42
    ),

    "XGBoost": XGBClassifier(
        n_estimators=80,
        learning_rate=0.1,
        max_depth=4,
        subsample=0.8,
    )
}

```

```

        colsample_bytree=0.8,
        eval_metric="logloss"
    )
}

# [2] RUN ALL MODELS
results = []
roc_data = {}

for name, model in models.items():
    clf = Pipeline(steps=[
        ("preprocessor", preprocessor),
        ("model", model)
    ])

    clf.fit(X_train, y_train)

    y_pred = clf.predict(X_test)
    y_prob = clf.predict_proba(X_test)[:, 1]

    results.append({
        "Model": name,
        "Accuracy": accuracy_score(y_test, y_pred),
        "Precision": precision_score(y_test, y_pred),
        "Recall": recall_score(y_test, y_pred),
        "F1-Score": f1_score(y_test, y_pred),
        "ROC-AUC": roc_auc_score(y_test, y_prob)
    })

    fpr, tpr, _ = roc_curve(y_test, y_prob)
    roc_auc = auc(fpr, tpr)
    roc_data[name] = (fpr, tpr, roc_auc)

results_df = pd.DataFrame(results)
print(results_df)

best = results_df.sort_values("ROC-AUC", ascending=False).iloc[0]
print("\nBest Model:\n", best)
best = results_df.sort_values("ROC-AUC", ascending=False).iloc[0]
print("Best Model:\n", best)
# [5] PLOT ROC CURVES
plt.figure(figsize=(10, 8))

```


```

for name, (fpr, tpr, roc_auc) in roc_data.items():
    plt.plot(fpr, tpr, lw=2, label=f"{name} (AUC = {roc_auc:.3f})")

# Baseline (no skill classifier)
plt.plot([0, 1], [0, 1], linestyle="--", color="gray")

plt.title("ROC Curve for All Models", fontsize=16)
plt.xlabel("False Positive Rate", fontsize=14)
plt.ylabel("True Positive Rate", fontsize=14)
plt.legend(loc="lower right")
plt.grid(True)
plt.show()

```



**Fig-1 Roc Curve of three models**

```
for name, model in models.items():
    print("\n" + "="*70)
    print(f"CONFUSION MATRIX & REPORT FOR: {name}")
    print("="*70)

    clf = Pipeline(steps=[
        ("preprocessor", preprocessor),
        ("model", model)
    ])
    clf.fit(X_train, y_train)

    y_pred = clf.predict(X_test)

    # Classification report
    print("\nClassification Report:")
    print(classification_report(y_test, y_pred))

    # Confusion matrix
    cm = confusion_matrix(y_test, y_pred)

    plt.figure(figsize=(6, 4))
    sns.heatmap(cm, annot=True, fmt="d", cmap="Blues")
    plt.title(f"Confusion Matrix - {name}")
    plt.xlabel("Predicted Label")
    plt.ylabel("True Label")
    plt.show()
```

Classification Report:


|              | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0            | 0.88      | 1.00   | 0.94     | 45139   |
| 1            | 0.00      | 0.00   | 0.00     | 5931    |
| accuracy     |           |        | 0.88     | 51070   |
| macro avg    | 0.44      | 0.50   | 0.47     | 51070   |
| weighted avg | 0.78      | 0.88   | 0.83     | 51070   |

=====  
CONFUSION MATRIX & REPORT FOR: Logistic Regression  
=====

Classification Report:

|              | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0            | 0.89      | 1.00   | 0.94     | 45139   |
| 1            | 0.60      | 0.03   | 0.06     | 5931    |
| accuracy     |           |        | 0.89     | 51070   |
| macro avg    | 0.75      | 0.51   | 0.50     | 51070   |
| weighted avg | 0.85      | 0.89   | 0.84     | 51070   |

Confusion Matrix — Logistic Regression




=====  
CONFUSION MATRIX & REPORT FOR: Random Forest  
=====

**Classification Report:**

|              | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0            | 0.88      | 1.00   | 0.94     | 45139   |
| 1            | 0.00      | 0.00   | 0.00     | 5931    |
| accuracy     |           |        | 0.88     | 51070   |
| macro avg    | 0.44      | 0.50   | 0.47     | 51070   |
| weighted avg | 0.78      | 0.88   | 0.83     | 51070   |

Confusion Matrix — Random Forest



---

**CONFUSION MATRIX & REPORT FOR: XGBoost**

---

Classification Report:

|              | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0            | 0.89      | 1.00   | 0.94     | 45139   |
| 1            | 0.63      | 0.06   | 0.11     | 5931    |
| accuracy     |           |        | 0.89     | 51070   |
| macro avg    | 0.76      | 0.53   | 0.52     | 51070   |
| weighted avg | 0.86      | 0.89   | 0.84     | 51070   |

Confusion Matrix — XGBoost

